

# Final\_AAI\_500\_Final\_Project\_Group3\_LogisticRegression

October 11, 2024

```
[1]: # Import necessary libraries - common libraries include pandas, numpy,
      ↪matplotlib, and sklearn
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import roc_curve, auc, roc_auc_score
from matplotlib import gridspec
import math
import scipy.stats
from scipy.stats import dgamma
from sklearn.linear_model import LinearRegression
from sklearn.datasets import fetch_california_housing
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix
from sklearn.preprocessing import StandardScaler
from imblearn.over_sampling import SMOTE
from sklearn.model_selection import GridSearchCV
import statsmodels.api as sm
from scipy.stats import skew, kurtosis
```

```
[2]: # Read in the Carbon West data file
spy = pd.read_csv('spy_Original.csv')

# View the first few rows of the dataset
spy.head()
```

```
[2]:
```

|   | Date      | Open      | High      | Low       | Close     | Volume  | Day | \ |
|---|-----------|-----------|-----------|-----------|-----------|---------|-----|---|
| 0 | 1/29/1993 | 24.701669 | 24.701669 | 24.578775 | 24.684113 | 1003200 | 29  |   |
| 1 | 2/1/1993  | 24.701643 | 24.859650 | 24.701643 | 24.859650 | 480500  | 1   |   |
| 2 | 2/2/1993  | 24.842113 | 24.929895 | 24.789444 | 24.912338 | 201300  | 2   |   |
| 3 | 2/3/1993  | 24.947451 | 25.193238 | 24.929894 | 25.175682 | 529400  | 3   |   |
| 4 | 2/4/1993  | 25.263461 | 25.333686 | 24.982561 | 25.281017 | 531500  | 4   |   |

|   | Weekday | Week | Month | Year |
|---|---------|------|-------|------|
| 0 | 4       | 4    | 1     | 1993 |
| 1 | 0       | 5    | 2     | 1993 |

|   |   |   |   |      |
|---|---|---|---|------|
| 2 | 1 | 5 | 2 | 1993 |
| 3 | 2 | 5 | 2 | 1993 |
| 4 | 3 | 5 | 2 | 1993 |

```
[3]: spy.tail()
```

```
[3]:
```

|      | Date      | Open       | High       | Low        | Close      | Volume \ |
|------|-----------|------------|------------|------------|------------|----------|
| 7949 | 8/26/2024 | 563.179993 | 563.909973 | 559.049988 | 560.789978 | 35788600 |
| 7950 | 8/27/2024 | 559.489990 | 562.059998 | 558.320007 | 561.559998 | 32693900 |
| 7951 | 8/28/2024 | 561.210022 | 561.650024 | 555.039978 | 558.299988 | 41066000 |
| 7952 | 8/29/2024 | 560.309998 | 563.679993 | 557.179993 | 558.349976 | 38715200 |
| 7953 | 8/30/2024 | 560.770019 | 564.200012 | 557.140015 | 563.679993 | 62700100 |

|      | Day | Weekday | Week | Month | Year |
|------|-----|---------|------|-------|------|
| 7949 | 26  | 0       | 35   | 8     | 2024 |
| 7950 | 27  | 1       | 35   | 8     | 2024 |
| 7951 | 28  | 2       | 35   | 8     | 2024 |
| 7952 | 29  | 3       | 35   | 8     | 2024 |
| 7953 | 30  | 4       | 35   | 8     | 2024 |

```
[4]: # check for null values and data type
spy.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7954 entries, 0 to 7953
Data columns (total 11 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Date        7954 non-null   object
1   Open        7954 non-null   float64
2   High        7954 non-null   float64
3   Low         7954 non-null   float64
4   Close       7954 non-null   float64
5   Volume      7954 non-null   int64
6   Day         7954 non-null   int64
7   Weekday     7954 non-null   int64
8   Week        7954 non-null   int64
9   Month       7954 non-null   int64
10  Year        7954 non-null   int64
dtypes: float64(4), int64(6), object(1)
memory usage: 683.7+ KB
```

```
[5]: # Convert the 'Date' column to a datetime format for proper date handling
spy['Date'] = pd.to_datetime(spy['Date'])

# Calculate the average 'Volume' prior to 2010 and post 2010
avg_volume_pre_2010 = spy[spy['Date'] < '2010-01-01']['Volume'].mean()
avg_volume_post_2010 = spy[spy['Date'] >= '2010-01-01']['Volume'].mean()
```

```

print("SPY AVG Volume Pior to 2010")
print(f"{avg_volume_pre_2010:.2f}")
print(' ')
print("SPY AVG Volume Post 2010")
print(f"{avg_volume_post_2010:.2f}")
print(' ')
print("SPY Volume Change Magnitude")
print(f"{avg_volume_post_2010 / avg_volume_pre_2010:.2f}")

```

SPY AVG Volume Pior to 2010  
58690971.58

SPY AVG Volume Post 2010  
113270950.27


SPY Volume Change Magnitude  
1.93

```

[6]: # set Date as the index
spy['Date_Index'] = pd.to_datetime(spy['Date']) # Convert to datetime if not_
      ↪already
spy.set_index('Date_Index', inplace=True) # Set 'Date' as the index
# Calculate the moving average
spy['SMA_100'] = spy['Close'].rolling(window=100).mean()
spy['SMA_200'] = spy['Close'].rolling(window=200).mean()


# plot spy performance since 2014 to visually trends and possible imbalances
plt.figure(figsize=(12, 6))
plt.plot(spy['Close'], label='Close Price', color='green')
plt.plot(spy['SMA_100'], label='100-Day SMA', color='yellow')
plt.plot(spy['SMA_200'], label='200-Day SMA', color='pink')
plt.title('SPY Price with Moving Averages')
plt.xlabel('Daily Data since 1993')
plt.ylabel('Price')
plt.legend()
plt.grid()
plt.show()

```



```
[7]: # plot distribution of returns
spy['Returns'] = spy['Close'].pct_change()

plt.figure(figsize=(12, 6))
plt.hist(spy['Returns'].dropna(), bins=50, color='green', alpha=0.7)
plt.title('Histogram of SPY Daily Returns')
plt.xlabel('Return')
plt.ylabel('Frequency')
plt.grid()
plt.show()
```



```
[9]: # Calculate daily returns using the 'Close' price
spy['Daily_Return'] = spy['Close'].pct_change() * 100

# Create subsets for the past 10 years and past 5 years
end_date = spy['Date'].max()
five_years_ago = end_date - pd.DateOffset(years=5)
ten_years_ago = end_date - pd.DateOffset(years=10)

spy_data_10y = spy[spy['Date'] >= ten_years_ago]
spy_data_5y = spy[spy['Date'] >= five_years_ago]

# New variables for skewness and kurtosis calculations
skewness_entire_dataset = skew(spy['Daily_Return'].dropna())
kurtosis_entire_dataset = kurtosis(spy['Daily_Return'].dropna())

skewness_last_10_years = skew(spy_data_10y['Daily_Return'].dropna())
kurtosis_last_10_years = kurtosis(spy_data_10y['Daily_Return'].dropna())

skewness_last_5_years = skew(spy_data_5y['Daily_Return'].dropna())
kurtosis_last_5_years = kurtosis(spy_data_5y['Daily_Return'].dropna())

# Display the results
print("Skewness and Kurtosis of SPY Daily Returns:")
print(f"Entire Dataset: Skewness = {skewness_entire_dataset:.2f}, Kurtosis = {kurtosis_entire_dataset:.2f}")
print(f"Past 10 Years: Skewness = {skewness_last_10_years:.2f}, Kurtosis = {kurtosis_last_10_years:.2f}")
```

```
print(f"Past 5 Years: Skewness = {skewness_last_5_years:.2f}, Kurtosis = {kurtosis_last_5_years:.2f}")
```

Skewness and Kurtosis of SPY Daily Returns:

Entire Dataset: Skewness = -0.06, Kurtosis = 11.38

Past 10 Years: Skewness = -0.54, Kurtosis = 12.69

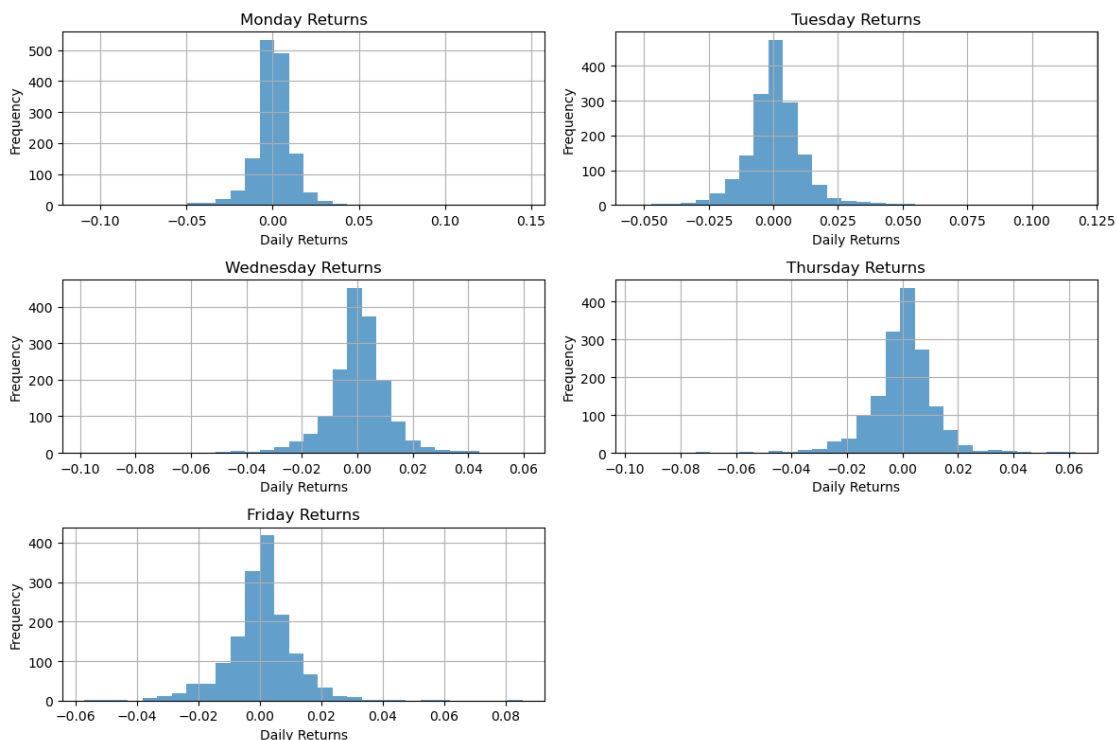
Past 5 Years: Skewness = -0.55, Kurtosis = 11.79

```
[10]: # Step 1: Calculate daily returns (if not already calculated)
spy['Daily_Returns'] = spy['Close'].pct_change()

# Step 2: Plot histograms for each day of the week using the existing 'Weekday'
# column (0 to 5)
# 0 = Monday, 5 = Friday
days = ['Monday', 'Tuesday', 'Wednesday', 'Thursday', 'Friday']

plt.figure(figsize=(12, 8))
for i, day in enumerate(days):
    plt.subplot(3, 2, i+1) # Create a subplot for each day
    spy[spy['Weekday'] == i]['Daily_Returns'].hist(bins=30, alpha=0.7)
    plt.title(f'{day} Returns')
    plt.xlabel('Daily Returns')
    plt.ylabel('Frequency')

plt.tight_layout() # Adjust layout to prevent overlap
plt.show()
```




```
[11]: # Step 1: Calculate daily returns (if not already calculated)
spy['Daily_Returns'] = spy['Close'].pct_change()

# Step 2: Categorize returns as positive (1) or negative (0)
spy['Positive_Return'] = spy['Daily_Returns'] > 0

# Step 3: Calculate the proportion of positive returns for each day of the week,
# using the existing 'Weekday' column
proportion_positive = spy.groupby('Weekday')['Positive_Return'].mean()

# Step 4: Plot the proportions
days = ['Monday', 'Tuesday', 'Wednesday', 'Thursday', 'Friday']

plt.figure(figsize=(8, 5))
plt.bar(days, proportion_positive, color='green', alpha=0.7)
plt.title('Proportion of Positive Returns by Day of the Week')
plt.xlabel('Day of the Week')
plt.ylabel('Proportion of Positive Returns')
plt.ylim(0, 1) # Proportions are between 0 and 1
plt.show()
```

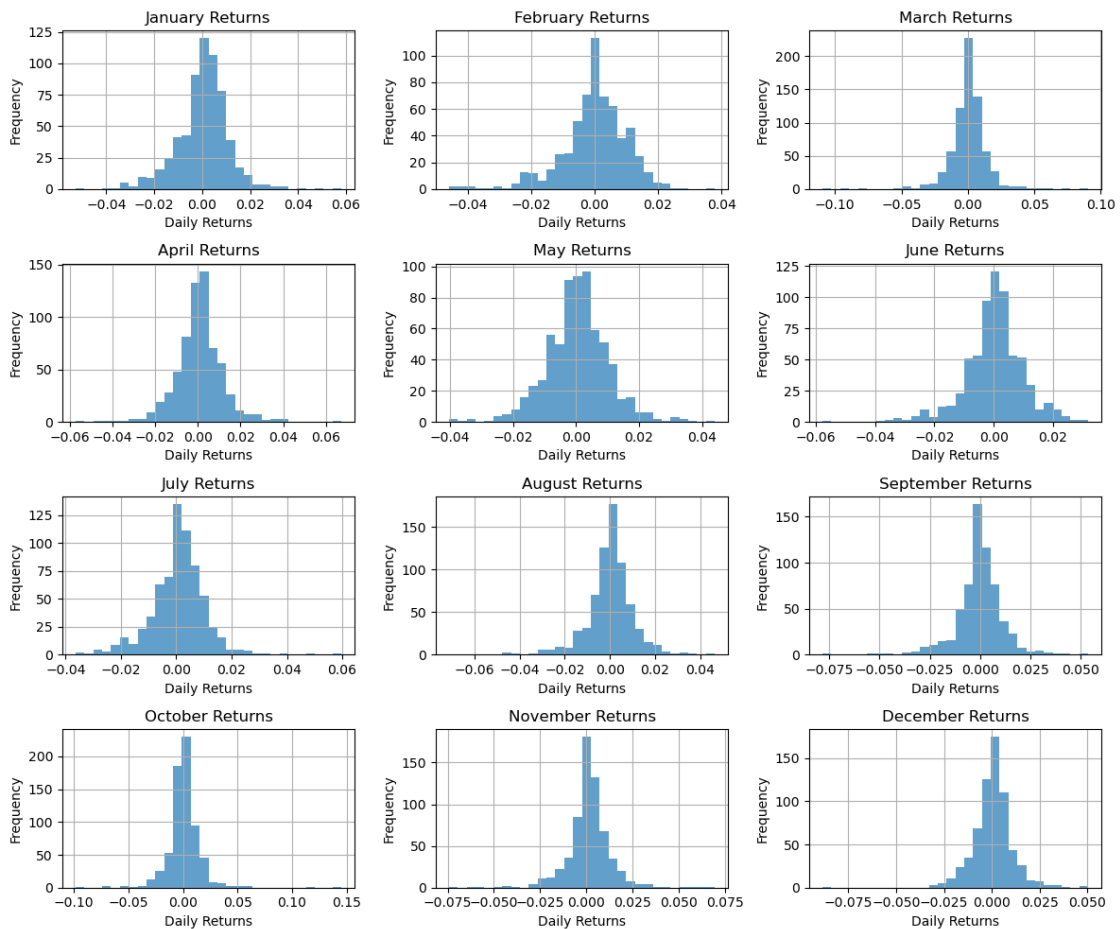


```
[12]: # Step 1: Calculate daily returns
spy['Daily_Returns'] = spy['Close'].pct_change()

# Step 3: Plot histograms of daily returns by month
months = ['January', 'February', 'March', 'April', 'May', 'June',
          'July', 'August', 'September', 'October', 'November', 'December']

plt.figure(figsize=(12, 10))
for i, month in enumerate(months):
    plt.subplot(4, 3, i+1) # Create a subplot for each month
    spy[spy['Month'] == (i+1)]['Daily_Returns'].hist(bins=30, alpha=0.7)
    plt.title(f'{month} Returns')
    plt.xlabel('Daily Returns')
    plt.ylabel('Frequency')

plt.tight_layout() # Adjust layout to prevent overlap
plt.show()
```






```
[13]: # Step 4: Categorize returns as positive (1) or negative (0)
spy['Positive_Return'] = spy['Daily_Returns'] > 0

# Step 5: Calculate the proportion of positive returns for each month
proportion_positive_month = spy.groupby('Month')['Positive_Return'].mean()

# Step 6: Plot the proportion of positive returns by month
plt.figure(figsize=(10, 6))
plt.bar(months, proportion_positive_month, color='blue', alpha=0.7)
plt.title('Proportion of Positive Returns by Month')
plt.xlabel('Month')
plt.ylabel('Proportion of Positive Returns')
plt.ylim(0, 1) # Proportions are between 0 and 1
plt.xticks(rotation=45) # Rotate x-axis labels for better readability
plt.show()
```



```
[19]: # Step 1: Create the binary target variable (1 if Close > Open, else 0)
spy['Outcome'] = (spy['Close'] > spy['Open']).astype(int)

# Step 2: Choose features and target variable
X = spy[['Open', 'High', 'Low', 'Volume']]
y = spy['Outcome']
```

```

# Step 3: Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    random_state=42)

# Step 4: Fit the logistic regression model
log_reg = LogisticRegression(max_iter=1000)
log_reg.fit(X_train, y_train)

# Step 5: Make predictions on the test set
y_pred = log_reg.predict(X_test)

# Evaluate model performance
accuracy = accuracy_score(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)

print(f"Logistic Regression Accuracy: {accuracy:.4f}")
print(f"Logistic Regression Confusion Matrix:")
print(conf_matrix)

```

```

Logistic Regression Accuracy: 0.7706
Logistic Regression Confusion Matrix:
[[456 284]
 [ 81 770]]

```

```

[29]: # Step 1: Create the binary target variable (1 if Close > Open, else 0)
spy['Outcome'] = (spy['Close'] > spy['Open']).astype(int)

# Create Exponential Moving Averages (EMA)
spy['EMA_5'] = spy['Close'].ewm(span=5, adjust=False).mean()
spy['EMA_10'] = spy['Close'].ewm(span=10, adjust=False).mean()
spy['EMA_20'] = spy['Close'].ewm(span=20, adjust=False).mean()

# Create rolling standard deviation for volatility
spy['Volatility_5'] = spy['Close'].rolling(window=5).std()
spy['Volatility_10'] = spy['Close'].rolling(window=10).std()

# Calculate Volume moving average
spy['Vol_MA_5'] = spy['Volume'].rolling(window=5).mean()
spy['Vol_MA_10'] = spy['Volume'].rolling(window=10).mean()

# Calculate Day of the Month
spy['Day_of_Month'] = pd.to_datetime(spy['Date']).dt.day

# Starting Features
spy['Range'] = spy['High'] - spy['Low'] # Add the daily range
spy['Prev_Close'] = spy['Close'].shift(1) # Previous day's close

```

```

spy['Prev_Open'] = spy['Open'].shift(1) # next day's open
spy['Intra_Move'] = spy['Prev_Close'] - spy['Prev_Open'] # get the previous day's
    ↳ intraday move
spy['Prev_Volume'] = spy['Volume'].shift(1) # Previous day's volume
spy['Prev_Return'] = (spy['Prev_Close'] - spy['Prev_Open']) / spy['Prev_Open']
    ↳ # Previous day's Intraday return
spy['Day_of_Week'] = pd.to_datetime(spy['Date']).dt.dayofweek # Add the
    ↳ day-of-week as a feature (0 for Monday, 4 for Friday, etc.)
spy['Month'] = pd.to_datetime(spy['Date']).dt.month # Extract the month from
    ↳ the Date column and add it as a feature

# Check the return from 2 days ago
spy['Lag_Return_2'] = spy['Prev_Return'].shift(1)

# Fill any missing values from shifting
spy.fillna(0, inplace=True)

# Define Features to use in model
X = spy[['Open', 'High', 'Low', 'Prev_Close', 'Prev_Volume', 'Prev_Return',
        'Day_of_Week', 'Intra_Move', 'EMA_5', 'EMA_10', 'EMA_20']]
y = spy['Outcome']

# Step 2: Feature Scaling
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Step 3: Handle class imbalance using SMOTE (Synthetic Minority Oversampling
    ↳ Technique)
# This allows us to handle imbalanced datasets which SPY is because over time
    ↳ it tends to increase
# SMOTE generates synthetic samples of the minority class to balance the dataset
smote = SMOTE(random_state=42)
# Fit the SMOTE algorithm to data set and resamples x_scaled and y to make both
    ↳ classes balanced
X_resampled, y_resampled = smote.fit_resample(X_scaled, y)

# Step 4: Split the resampled data into training and testing sets
# split data by 80/20, train on 80% and test on 20%
X_train, X_test, y_train, y_test = train_test_split(X_resampled, y_resampled,
    ↳ test_size=0.2, random_state=42)

# Step 5: Train the logistic regression model with class weights to balance
    ↳ classes
param_grid = {'C': [0.01, 0.1, 1, 10, 100]}
log_reg = LogisticRegression(max_iter=1000, class_weight='balanced')
grid_search = GridSearchCV(log_reg, param_grid, cv=5)

```

```

# Train the model with grid search
# log_reg.fit(X_train, y_train)
grid_search.fit(X_train, y_train)

# Best model from grid search
best_log_reg = grid_search.best_estimator_

# Step 6: Make predictions on the test set
#y_pred = log_reg.predict(X_test)
# Make predictions with the best model
y_pred = best_log_reg.predict(X_test)

# Evaluate model performance
accuracy = accuracy_score(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)

print(f"Logistic Regression Accuracy: {accuracy:.4f}")
print(f"Logistic Regression Confusion Matrix:")
print(conf_matrix)

```

```

Logistic Regression Accuracy: 0.8433
Logistic Regression Confusion Matrix:
[[711 115]
 [146 694]]


```

```

[30]: # Run correlation on the features
correlation_matrix = X.corr()

# Create a heatmap of the correlation matrix using Seaborn
plt.figure(figsize=(10, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt=".2f",
            linewidths=0.5)
plt.title('Logistic Regression Feature Correlation Matrix Heatmap')
plt.show()

```



```
[25]: # Extract the values
TN, FP, FN, TP = conf_matrix[0][0], conf_matrix[0][1], conf_matrix[1][0],
↪conf_matrix[1][1]

# Calculating Precision, Recall, and F1-score manually
precision = TP / (TP + FP)
recall = TP / (TP + FN)
f1 = 2 * (precision * recall) / (precision + recall)

print(f"Precision: {precision:.4f}")
print(f"Recall: {recall:.4f}")
print(f"F1-score: {f1:.4f}")
```

```
Precision: 0.8578
Recall: 0.8262
F1-score: 0.8417
```

```
[26]: # Add a constant for the intercept term
X_resampled_with_const = sm.add_constant(X_resampled)

# Perform logistic regression using statsmodels to get the p-values for the
↳Wald test
logit_model = sm.Logit(y_resampled, X_resampled_with_const)
result = logit_model.fit()

# Get the summary of the model which includes the Wald test results (p-values)
summary = result.summary()
summary
```

Optimization terminated successfully.

Current function value: 0.359599

Iterations 10

[26]:

|                         |                  |                          |         |
|-------------------------|------------------|--------------------------|---------|
| <b>Dep. Variable:</b>   | Outcome          | <b>No. Observations:</b> | 8328    |
| <b>Model:</b>           | Logit            | <b>Df Residuals:</b>     | 8316    |
| <b>Method:</b>          | MLE              | <b>Df Model:</b>         | 11      |
| <b>Date:</b>            | Fri, 11 Oct 2024 | <b>Pseudo R-squ.:</b>    | 0.4812  |
| <b>Time:</b>            | 09:48:12         | <b>Log-Likelihood:</b>   | -2994.7 |
| <b>converged:</b>       | True             | <b>LL-Null:</b>          | -5772.5 |
| <b>Covariance Type:</b> | nonrobust        | <b>LLR p-value:</b>      | 0.000   |

|              | coef      | std err | z       | P>  z | [0.025   | 0.975]   |
|--------------|-----------|---------|---------|-------|----------|----------|
| <b>const</b> | 0.0261    | 0.043   | 0.608   | 0.543 | -0.058   | 0.110    |
| <b>x1</b>    | -666.7706 | 17.673  | -37.729 | 0.000 | -701.409 | -632.133 |
| <b>x2</b>    | 236.1724  | 10.848  | 21.772  | 0.000 | 214.911  | 257.434  |
| <b>x3</b>    | 280.3360  | 10.062  | 27.860  | 0.000 | 260.614  | 300.058  |
| <b>x4</b>    | -261.2612 | 9.800   | -26.659 | 0.000 | -280.469 | -242.054 |
| <b>x5</b>    | 0.0114    | 0.039   | 0.296   | 0.768 | -0.064   | 0.087    |
| <b>x6</b>    | -0.0645   | 0.058   | -1.119  | 0.263 | -0.177   | 0.048    |
| <b>x7</b>    | 0.0105    | 0.032   | 0.330   | 0.742 | -0.052   | 0.073    |
| <b>x8</b>    | 1.0819    | 0.098   | 11.002  | 0.000 | 0.889    | 1.275    |
| <b>x9</b>    | 937.3592  | 34.062  | 27.520  | 0.000 | 870.600  | 1004.118 |
| <b>x10</b>   | -726.8857 | 29.232  | -24.866 | 0.000 | -784.180 | -669.591 |
| <b>x11</b>   | 201.1291  | 10.008  | 20.096  | 0.000 | 181.513  | 220.745  |

```
[31]: # Visualize confusion matrix from Logistic Regression
# Step 1: Confusion Matrix Heatmap
plt.figure(figsize=(8, 6))
sns.heatmap(conf_matrix, annot=True, fmt="d", cmap="Blues", cbar=False)
plt.title("Logistic Regression Results Confusion Matrix")
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.show()

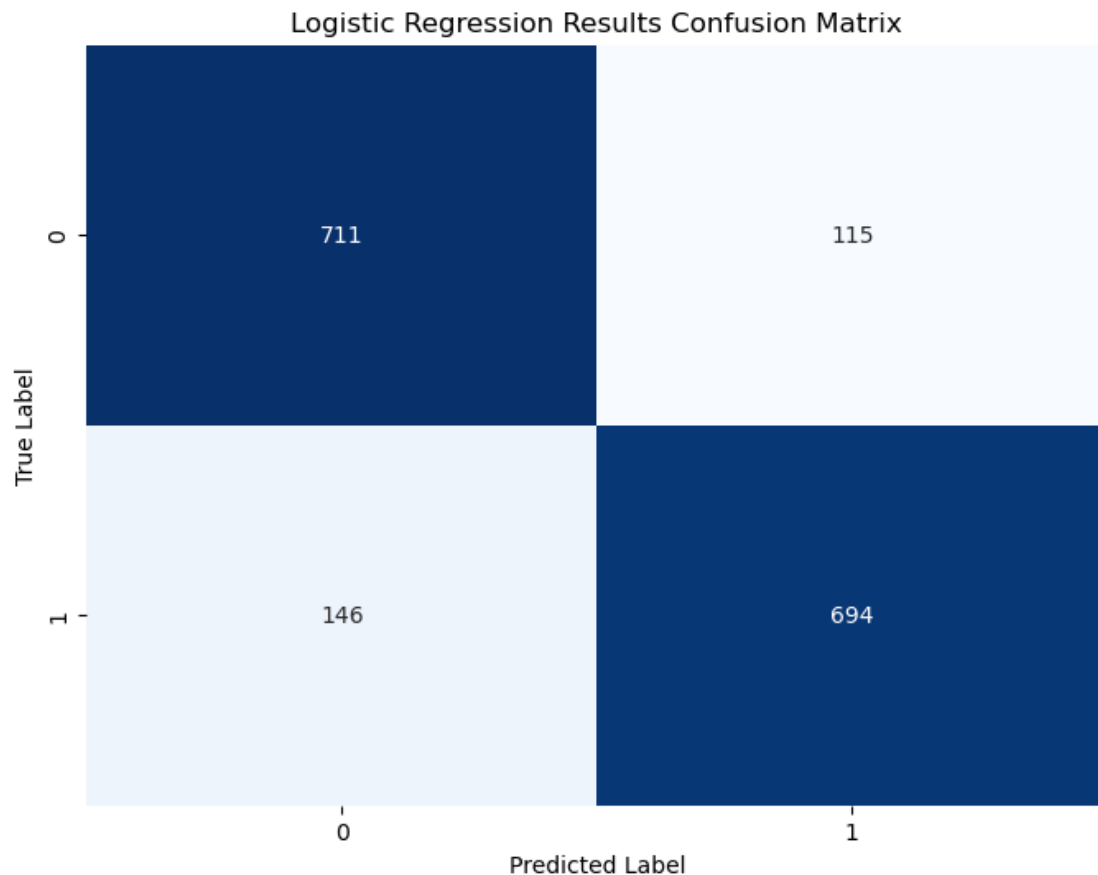
# Step 2: Plot ROC Curve
```


```

y_prob = best_log_reg.predict_proba(X_test)[: , 1] # Probabilities for the
↪positive class
fpr, tpr, thresholds = roc_curve(y_test, y_prob)
roc_auc = auc(fpr, tpr)

plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, label=f"ROC Curve (AUC = {roc_auc:.2f})", color="darkorange")
plt.plot([0, 1], [0, 1], color="navy", linestyle="--")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("SPY Logistic Regression Model (ROC) Curve")
plt.legend(loc="lower right")
plt.grid()
plt.show()

```





```
[17]: from sklearn.model_selection import cross_val_score
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.neural_network import MLPClassifier
from sklearn.ensemble import GradientBoostingClassifier
import numpy as np

# Define the models to compare
models = {
    "Logistic Regression": LogisticRegression(max_iter=1000,
↪class_weight='balanced'),
    "Decision Tree": DecisionTreeClassifier(),
    "Grandient Boosting": GradientBoostingClassifier(),
}
```



```

# Perform 5-fold cross-validation for each model
results = {}
for model_name, model in models.items():
    # Cross-validate the model using accuracy as the metric
    scores = cross_val_score(model, X_scaled, y, cv=5, scoring='accuracy')
    results[model_name] = np.mean(scores), np.std(scores)
    print(f"{model_name}: Mean Accuracy = {np.mean(scores):.4f}, Std = {np.
    ↪std(scores):.4f}")

# Print the summary of results
results

```

Logistic Regression: Mean Accuracy = 0.5591, Std = 0.0589  
 Decision Tree: Mean Accuracy = 0.4689, Std = 0.0424  
 Gradient Boosting: Mean Accuracy = 0.4581, Std = 0.0462

```

[17]: {'Logistic Regression': (0.5590862912056418, 0.05888021479372918),
      'Decision Tree': (0.4689495550838245, 0.04235667607287223),
      'Gradient Boosting': (0.4581389814562259, 0.046163505372091296)}

```

```

[19]: # Define a function to randomly sample values from the dataset and use them in
    ↪the prediction model
def random_sample_prediction(spy, model, scaler):
    # Step 1: Randomly sample one row from the dataset
    random_sample = spy.sample(1)

    # Step 2: Extract features required for the prediction
    random_sample_data = random_sample[['Open', 'High', 'Low', 'Prev_Close',
    ↪'Prev_Volume', 'Prev_Return',
    'Day_of_Week', 'Intra_Move', 'EMA_5',
    ↪'EMA_10', 'EMA_20']]

    # Step 3: Scale the features using the existing scaler
    random_sample_scaled = scaler.transform(random_sample_data)

    # Step 4: Make a prediction using the trained logistic regression model
    prediction = model.predict(random_sample_scaled)
    prediction_proba = model.predict_proba(random_sample_scaled)

    # Step 5: Return the prediction and probabilities
    result = {
        'Sampled Data': random_sample_data,
        'Prediction': 'Close > Open' if prediction[0] == 1 else 'Close <= Open',
        'Probability (Close > Open)': prediction_proba[0][1],
        'Probability (Close <= Open)': prediction_proba[0][0]
    }

```

```

    return result

# Run the random sampling prediction using the current dataset, trained model, and scaler
random_sample_result = random_sample_prediction(spy, best_log_reg, scaler)
random_sample_result

```

```

[19]: {'Sampled Data':          Open          High          Low  Prev_Close
Prev_Volume \
5781  167.813479  168.167319  162.574979  167.131729  172330500.0

          Prev_Return  Day_of_Week  Intra_Move          EMA_5          EMA_10          EMA_20
5781      -0.000826           2    -0.138086  166.317094  168.978531  171.81041
,
'Prediction': 'Close <= Open',
'Probability (Close > Open)': 0.0059758378617140685,
'Probability (Close <= Open)': 0.9940241621382859}

```

```

[32]: import pandas as pd
import numpy as np
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix

def process_and_train_model(file_path):
    # Load the data from the file path
    spy = pd.read_csv(file_path)

    # Step 1: Create the binary target variable (1 if Close > Open, else 0)
    spy['Outcome'] = (spy['Close'] > spy['Open']).astype(int)

    # Create Exponential Moving Averages (EMA)
    spy['EMA_5'] = spy['Close'].ewm(span=5, adjust=False).mean()
    spy['EMA_10'] = spy['Close'].ewm(span=10, adjust=False).mean()
    spy['EMA_20'] = spy['Close'].ewm(span=20, adjust=False).mean()

    # Create rolling standard deviation for volatility
    spy['Volatility_5'] = spy['Close'].rolling(window=5).std()
    spy['Volatility_10'] = spy['Close'].rolling(window=10).std()

    # Calculate Volume moving average
    spy['Vol_MA_5'] = spy['Volume'].rolling(window=5).mean()
    spy['Vol_MA_10'] = spy['Volume'].rolling(window=10).mean()

    # Calculate Day of the Month

```

```

spy['Day_of_Month'] = pd.to_datetime(spy['Date']).dt.day

# Feature Engineering
spy['Range'] = spy['High'] - spy['Low'] # Add the daily range
spy['Prev_Close'] = spy['Close'].shift(1) # Previous day's close
spy['Prev_Open'] = spy['Open'].shift(1) # Next day's open
spy['Intra_Move'] = spy['Prev_Close'] - spy['Prev_Open'] # Get the
↳previous day intraday move
spy['Prev_Volume'] = spy['Volume'].shift(1) # Previous day's volume
spy['Prev_Return'] = (spy['Prev_Close'] - spy['Prev_Open']) /
↳spy['Prev_Open'] # Previous day's Intraday return
spy['Day_of_Week'] = pd.to_datetime(spy['Date']).dt.dayofweek # Add the
↳day-of-week as a feature
spy['Month'] = pd.to_datetime(spy['Date']).dt.month # Extract the month
↳from the Date column
# Check the return from 2 days ago
spy['Lag_Return_2'] = spy['Prev_Return'].shift(1)

# Fill any missing values from shifting
spy.fillna(0, inplace=True)

# Define Features to use in model
X = spy[['Open', 'High', 'Low', 'Prev_Close', 'Prev_Volume', 'Prev_Return',
        'Day_of_Week', 'Intra_Move', 'EMA_5', 'EMA_10', 'EMA_20']]
y = spy['Outcome']

# Step 2: Feature Scaling
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Step 3: Handle class imbalance using the built-in balanced class weights
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y,
↳test_size=0.2, random_state=42)

# Step 4: Train the logistic regression model with grid search
param_grid = {'C': [0.01, 0.1, 1, 10, 100]}
log_reg = LogisticRegression(max_iter=1000, class_weight='balanced')
grid_search = GridSearchCV(log_reg, param_grid, cv=5)
grid_search.fit(X_train, y_train)

# Best model from grid search
best_log_reg = grid_search.best_estimator_

# Step 5: Make predictions on the test set
y_pred = best_log_reg.predict(X_test)

# Step 6: Evaluate model performance

```

```

accuracy = accuracy_score(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)

# Return the accuracy, confusion matrix, and trained model
return accuracy, conf_matrix, best_log_reg, scaler

```

```

[37]: import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import roc_curve, auc

# Call the function with your data file
file_path = 'GOOGL.csv'
accuracy, conf_matrix, trained_model, scaler = _
    ↪ process_and_train_model(file_path)

# Display accuracy and confusion matrix
print(f"Accuracy: {accuracy:.4f}")
print("Confusion Matrix:")
print(conf_matrix)

# Visualize confusion matrix from Logistic Regression
# Step 1: Confusion Matrix Heatmap
plt.figure(figsize=(8, 6))
sns.heatmap(conf_matrix, annot=True, fmt="d", cmap="Blues", cbar=False)
plt.title("Confusion Matrix")
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.show()


# Step 2: Plot ROC Curve
y_prob = best_log_reg.predict_proba(X_test)[:, 1] # Probabilities for the _
    ↪ positive class
fpr, tpr, thresholds = roc_curve(y_test, y_prob)
roc_auc = auc(fpr, tpr)

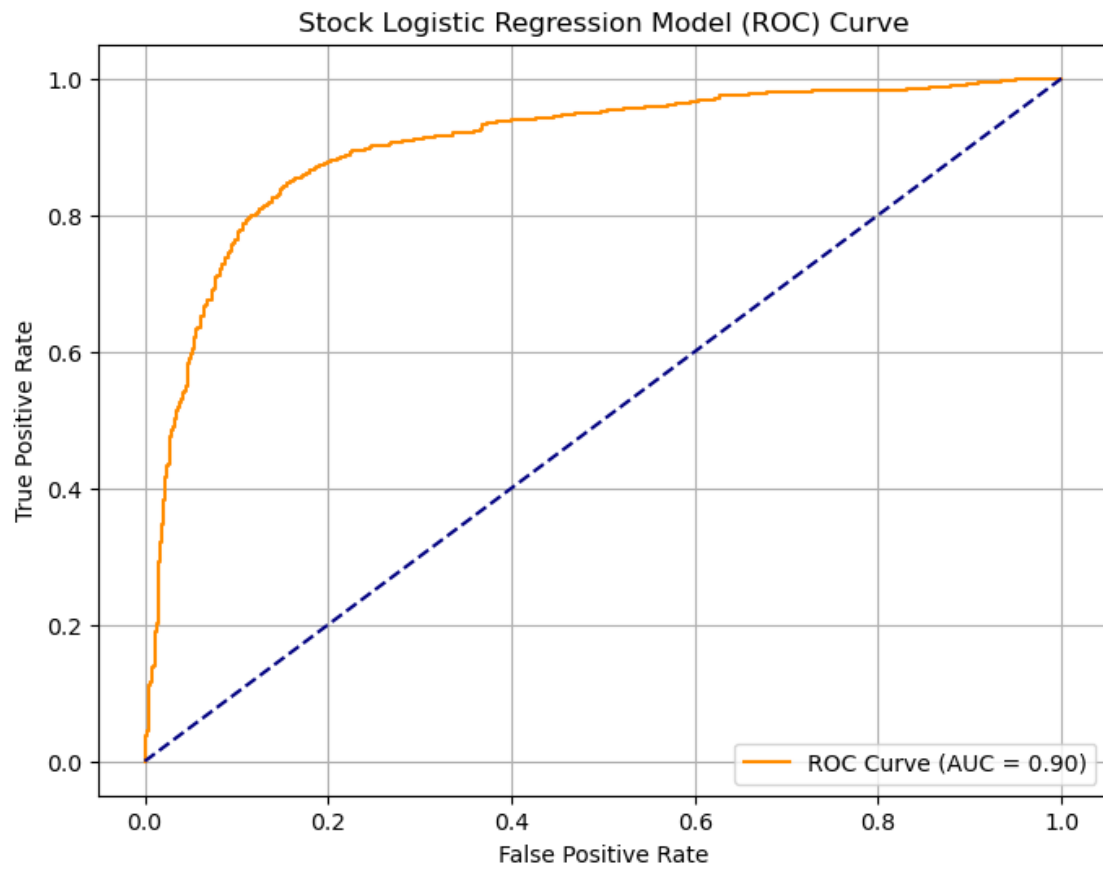
plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, label=f"ROC Curve (AUC = {roc_auc:.2f})", color="darkorange")
plt.plot([0, 1], [0, 1], color="navy", linestyle="--")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("Stock Logistic Regression Model (ROC) Curve")
plt.legend(loc="lower right")
plt.grid()
plt.show()

```

Accuracy: 0.8444  
 Confusion Matrix:  
 [[373 67]

```
[ 71 376]]
```





[ ]: